

Mobile “Offline-Only” BBQ Competition Server — Product Specification (v0.1)

Purpose: Provide a self-contained, event-ready, offline web platform to run BBQ competitions end-to-end (setup → registration → turn-in → judging → scoring → reporting) **without relying on external Wi-Fi or Internet**. All participants operate **only** on the event network, and all data is stored locally with **near-zero tolerance for downtime or corruption**.

1) Non-Negotiable Requirements

1.1 Availability & Data Integrity

- **Database must not go down** during an event. A DB outage is an event failure.
 - **Database must not get corrupted** (power loss, storage failure, process crash must not cause irrecoverable data loss).
 - System must tolerate:
 - Single node failure (server/worker Pi failure)
 - Single SSD failure (in the database layer)
 - Power interruption long enough to complete safe shutdown (UPS-driven)

Definition of success: The event can complete and produce final rankings and team reports even if one compute node fails mid-event.

1.2 Offline-Only Network / “Walled Garden”

- All client devices must use the event network and reach only:
 - the load balancer
 - the application services
 - DNS/captive portal services (if enabled)
- Internet access is blocked by default (optionally allowed for admins only).

1.3 Event Simplicity & Operator UX

- Host arrives, flips **one master power switch**.
 - Cluster self-boots into “Ready” or “Not Ready” state.
 - Physical indicators:
 - Per-node LED (Green = healthy, Red = unhealthy)
 - Master LED (Green = cluster ready, Red = cluster not ready)
 - Minimal steps to start an event:
 - “Create Event”

- “Configure Categories / Rules / Weighting”
 - “Open Registration”
 - “Open Turn-in”
 - “Start Judging”
 - “Finalize / Publish Results”
 - “Generate Reports”
-

2) System Overview

2.1 Roles

- **Host/Admin (staff):** Configure competition, manage check-in, validate turn-ins, orchestrate judging flow, finalize results, print/export reports.
 - **Judge:** Uses personal device (phone/tablet) to score assigned samples. Must be guided; should not choose categories manually.
 - **Runner/Table Host (staff):** Physically distributes samples and coordinates “appearance” then “taste/texture” rounds per table.
 - **Contestant (optional access):** If enabled, can view their own reports after completion (or staff prints them).

2.2 Core Workflows

1. Event Setup

- Define categories (e.g., chicken, pulled pork, ribs, brisket)
 - Define grading criteria (appearance, taste, texture)
 - Define grading scale (7–10 with labels)
 - Define weighting (per criterion and per category)
 - Define number of judge tables and seats per table

2. Registration

- Team registers
 - System assigns unique Team ID
 - System produces scannable team code (Aztec/PDF417 recommended)
 - Team code is printed and given to team

3. Turn-in

- Each submission has a pre-printed identifier (submission code)
 - Staff scans submission code + scans team code to validate ownership
 - System records timestamp and marks submission “Received”

4. Judging

- Judges scan table QR to open table-specific URL
 - Judge enters seat number (1–6)

- Server assigns the judge's next sample, consistent with the table's distribution plan
- Flow:
 - **Appearance phase:** judges evaluate all assigned samples in numeric order (1..N at table)
 - **Taste/Texture phase:** judges evaluate in seat-specific "passing order"
- Comments allowed for each criterion

5. Scoring & Results

- System calculates per-team category score
 - Produces ranking best → worst for each category
 - Calculates overall winner using host-defined category weights
 - Produces team report including aggregates + all comments
-

3) Hardware Architecture (High Availability, Offline, Event-Proof)

3.1 Guiding Principle

Do not use Raspberry Pi as the Wi-Fi access point for 200 clients. Use proper AP hardware and keep Pis as compute/storage.

3.2 Recommended Physical Stack (Minimum "No DB Downtime")

This is the baseline that matches the "DB cannot go down" requirement.

A) Network / Edge

- 1× **Router/Firewall appliance** (small x86 or high-quality travel router with VLAN + firewall rules)
 - 2× **Wi-Fi Access Points** (capacity planning for 200 devices; dual-band, tuned for high client count)
 - Optional: managed switch (PoE if APs need it)

B) Compute Cluster

- 2× **Load Balancer Nodes**
 - Can be Raspberry Pi 4/5
 - Runs HAProxy/Nginx + Keepalived (Virtual IP failover)
- 2–3× **Application Worker Nodes (RPi 5 recommended)**
 - Runs Node.js API and serves the SPA static site
 - Horizontal scale: add more workers as needed

C) Database Layer (Hard Requirement)

- 2× **Database Nodes** (dedicated)

- Each with **USB3 SSD** (high endurance)
- Runs PostgreSQL in an HA configuration with synchronous replication
- 1x **Witness/Quorum Node** (lightweight)
 - Prevents split-brain
 - Can be a small Pi or run on one LB node if isolated appropriately

Total: 2 LB + 2 DB + 1 Witness + 2–3 App Workers = **7–8 devices**

(You can combine Witness onto an LB node if you must reduce count → **6–7 devices**, but separate is cleaner.)

3.3 Power & Corruption Protection

- **UPS is mandatory** for DB integrity:
 - UPS powers router/switch/AP + DB nodes + at least one LB and one App worker
 - UPS triggers clean shutdown on low battery
- **Storage:**
 - DB nodes use SSD only; **no microSD for DB data**
 - Enable filesystem and DB settings for durability (see §7.3)
- **Enclosure:**
 - Rugged case with airflow
 - Labeling for ports and node roles
 - Single master toggle switch controlling a power strip feeding all components

3.4 LED Health Indicator System

- Each node exposes a local health endpoint or exports status to the "Indicator Controller".
 - **Indicator Controller** (can be one LB node or a dedicated small Pi):
 - Polls:
 - LB active/standby status
 - App workers health (HTTP 200 + dependency checks)
 - DB cluster leader/replica health + replication lag
 - Disk SMART + filesystem read/write checks
 - UPS status (battery/line power)
 - Drives:
 - Per-node LED: Green/Red
 - Master LED:
 - Green if minimum quorum healthy:
 - (LB VIP is up)
 - (≥ 1 App worker healthy)
 - (DB leader healthy AND replication healthy OR failover-

- ready)
– Red otherwise
-

4) Network Design ("Only on Our Network")

4.1 SSID & Client Access

- 1 SSID for judges/participants (WPA2/WPA3 PSK)
 - Optional separate SSID for staff/admin devices

4.2 DNS & URL Experience

- Preferred: <https://event.local/> or <http://event.local/>
 - Local DNS server on router or cluster resolves `event.local` to the LB virtual IP
 - Optional captive portal to auto-open the judging web app upon connection

4.3 Firewall Rules

- Block outbound Internet for participant SSID VLAN
 - Allow participant VLAN → LB VIP ports only (80/443)
 - Staff VLAN can reach admin endpoints
 - Optional: block device-to-device traffic on participant VLAN (client isolation)

5) Software Platform Recommendations

5.1 Front-End

- **React** (recommended) or Angular
 - Must be **mobile-first** and work on:
 - iOS Safari
 - Android Chrome
 - Implement as **PWA** for:
 - fast reload
 - resilience to brief Wi-Fi drops
 - optional "Add to Home Screen"
 - Keep UI simple and "guided":
 - Judges never choose category; server assigns next step.

5.2 Back-End API

- **Node.js** with a structured framework:
 - **NestJS** recommended (strong modularity, validation, DI)
 - API patterns:

- REST for most operations
- Optional WebSockets/SSE for live status dashboards (host view)

5.3 Database

- **PostgreSQL** (primary choice)
 - HA Layer:
 - Patroni (or similar) for leader election and automated failover
 - Synchronous replication between DB nodes
 - Quorum/witness to prevent split-brain

5.4 Cache / Queue (Optional but Strongly Recommended)

- Redis for:
 - rate-limiting
 - short-lived session tokens
 - buffering telemetry/metrics
 - If you can accept "acknowledged then persisted" semantics:
 - message queue for smoothing spikes
 - BUT final persistence must be in Postgres
-

6) Data Model (Conceptual)

6.1 Entities

- **Event**
 - id, name, date, location, status (setup/registration/turnin/judging/ finalized)
- **Category**
 - id, event_id, name (chicken, ribs, etc.), enabled
- **Criteria**
 - appearance/taste/texture (baseline), configurable list
- **Scale**
 - values 7–10 with labels (poor/fair/good/excellent)
- **Weighting**
 - per-category criterion weights (e.g., taste 0.5, texture 0.3, appearance 0.2)
 - overall category weight for grand champion calc
- **Team**
 - id, event_id, team_number, team_name, metadata
- **TeamCode**
 - team_id, code_type (Aztec/PDF417), payload, printed_at
- **Submission**
 - id, event_id, category_id, submission_number, submission_code, status

- **TurnIn**
 - submission_id, scanned_team_id, verified (bool), timestamp, operator_id
- **JudgeTable**
 - id, event_id, table_number, seats (fixed at 6 in this spec)
- **JudgeSeat**
 - table_id, seat_number (1–6), active_session_token, judge_alias (optional)
- **AssignmentPlan**
 - category_id, table_id, list of submission_ids assigned to that table
- **Score**
 - submission_id, judge_seat_id, criterion_id, value (7–10), comment, timestamp
- **AuditLog**
 - append-only: who did what, when, from which device

6.2 Integrity Constraints

- One score per (submission_id, judge_seat_id, criterion_id, phase)
 - Turn-in must record both scanned submission + scanned team, and verify team ownership
 - Submission cannot be judged until status = "AssignedToTables" and judging phase is active
-

7) Judging Logic & Sequencing

7.1 Table Assignment

Inputs:

- number_of_judges = tables × 6
 - total_submissions per category
 - Output:
 - Each table receives $\lceil \text{total_submissions} / \text{tables} \rceil$ items (balanced as evenly as possible)
 - Randomization:
 - Randomize submissions to tables
 - Randomize initial order within each table list (but assign a stable numeric order for appearance phase)

7.2 Appearance Phase

- Judges evaluate appearance for **all submissions at their table**
 - Order: numerical (1..N) within that table's list
 - Each score:
 - value 7–10

- optional comment
- System enforces completion gating:
 - Taste/texture phase cannot start until appearance complete for the table (or per-judge; choose one policy and enforce consistently)

7.3 Taste/Texture Phase

- Judges evaluate taste and texture per submission, following the seat-specific “passing order”
 - The system must compute the next submission for each seat based on:
 - Table submission count N
 - Seat number s (1–6)
 - A defined deterministic permutation that matches physical passing behavior

Requirement: The seat ordering algorithm must reproduce the patterns described:

- Example for seat 1: 1, 7, 8, 9, 10, 11, 12, 13, 14, 15, 6, 5, 4, 3, 2
 - Example for seat 6: 6, 5, 4, 3, 2, 1, 15, 14, 13, 12, 11, 10, 9, 8, 7

Implementation approach:

- Represent table seats in a ring and define a “pass direction”
 - Model “batches” of 6 (each round of passing)
 - Generate a sequence per seat as a function of:
 - initial seat assignment
 - pass direction
 - serving entry point (seat 3 per your description)
 - Store the generated sequence per (table, category) so it is stable and auditable.

7.4 Comments

- Comments allowed for each criterion on each submission
 - Host reports must include all comments (with judge anonymization if desired)

8) Scoring & Reporting

8.1 Per-Submission Scoring

For each submission:

- Aggregate per criterion across 6 judges:
 - mean or trimmed mean (choose policy)
- Weighted criterion score:
 - `criterion_score × criterion_weight`

8.2 Per-Team Category Score

- Team score for category = submission score (assuming one submission per category)
 - If multiple submissions per category are possible, define policy (best-of, average, etc.)

8.3 Overall Score

- Overall team score = $\text{sum}(\text{category_score} \times \text{category_weight})$
 - Output:
 - category rankings
 - overall rankings

8.4 Team Report

Per team:

- Per category:
 - appearance avg + comments
 - taste avg + comments
 - texture avg + comments
 - overall category score
 - Overall placement (if finalized)

Exports:

- PDF report per team
 - CSV export of all scores (for audit/recovery)

9) Security & Anti-Tamper

9.1 Codes (Team & Submission)

- Use **Aztec or PDF417** codes for printed identifiers
 - Payload should be:
 - opaque token (UUID + signature)
 - not human-readable IDs
 - Include a signature/HMAC so codes can't be forged offline

9.2 Judge Access

- Table QR encodes:
 - table_id
 - signed token with expiry
- Judge enters seat number (1–6)
- Server issues seat session token:
 - short-lived access token

- renew silently
- Prevent cross-seat scoring:
 - seat token binds to seat_id; API rejects mismatch

9.3 Admin Access

- Admin UI available only on staff VLAN (recommended)
 - Two-level permissions:
 - “Operator” (turn-in, table management)
 - “Administrator” (setup, finalize, exports)
-

10) Reliability Engineering (“DB Cannot Go Down”)

10.1 Database HA Strategy

- Two Postgres nodes with **synchronous replication**
 - Leader commits only when replica confirms write (prevents data loss on leader failure)
 - Automated failover with Patroni (or equivalent)
 - Quorum/witness to avoid split-brain
 - Health checks enforced at LB and app tiers:
 - If DB is degraded, app shows “pause” state rather than accepting writes blindly

10.2 Storage Strategy

- Each DB node uses:
 - SSD with high endurance
 - Proper filesystem (ext4 recommended) with safe mount options
- Optional but strong:
 - Mirror SSDs (RAID1) per DB node if hardware supports it (often easier on x86 than Pi)

10.3 Power Strategy

- UPS required.
 - Auto graceful shutdown on low battery.
 - On boot:
 - DB cluster checks consistency
 - App stays in “Not Ready” until DB leader elected and replication healthy

10.4 Backup Strategy During Event

- Continuous:
 - write-ahead log archiving to the second DB node
- Periodic (e.g., every 2 minutes):

- logical snapshot export to a third storage location (USB SSD on controller)
 - End-of-round:
 - export CSV + PDF packs
 - Outcome:
 - Even in catastrophic failure, you have a near-real-time recovery point.
-

11) Implementation Plan (Phased Rollout)

Phase 0 — Prototype (Single Node)

- One machine with Postgres + Node + React
 - Implement:
 - event setup
 - registration + codes
 - turn-in verification
 - basic judging UI
 - scoring + reports

Phase 1 — Production MVP (No-Downtime DB Baseline)

- Introduce:
 - LB VIP failover
 - 2 app workers
 - 2 DB nodes synchronous replication
 - UPS + health indicator controller
- Add:
 - audit logs
 - admin/staff roles
 - exports (CSV/PDF)

Phase 2 — Event Operations Polish

- Captive portal / easy "connect & judge"
 - Host dashboards:
 - turn-in progress per category
 - judging completion by table/seat
 - "stuck judge" detection
 - Better randomization + repeatability (seeded randomness)

Phase 3 — Hardening

- Fault injection tests:
 - kill app node
 - kill DB leader
 - unplug SSD

- power-loss simulation on UPS
 - Load testing for:
 - 200 concurrent judges
 - 30–60 writes/sec bursts
-

12) Acceptance Criteria

Availability / Integrity

- If any single app worker dies, system continues without loss.
 - If DB leader dies, replica is promoted automatically and writes continue.
 - If power fails:
 - UPS keeps system alive long enough to shut down cleanly.
 - On restart, event data is intact and resumes.

Performance

- Supports:
 - 200 concurrent clients
 - typical 5–10 API requests per device per minute
 - burst handling during “save scores” moments

Operator UX

- “Power on → Ready” within a defined boot window (target: <5 minutes).
 - Master LED accurately signals readiness.
 - Judges can score with minimal steps:
 - scan → seat → score
-

13) Open Decisions (Project Manager Must Lock Early)

1. Final DB HA tooling: Patroni vs another approach.
 2. Trimmed mean vs average scoring policy; tie-break rules.
 3. Whether taste/texture scoring is per-submission sequential locking or allows out-of-order entry.
 4. Judge anonymity policy in team reports.
 5. Captive portal requirement vs simple DNS URL.
-

14) Deliverables Checklist (Engineering)

- Infrastructure-as-code for node provisioning
 - Docker images / compose files for:
 - LB + VIP
 - App worker
 - DB HA stack
 - Monitoring + LED controller
 - Full schema migrations + seed data

- Admin UI + Judge UI + Operator UI
 - PDF generation service and templates
 - Load tests + chaos tests
 - Event day runbook:
 - setup checklist
 - troubleshooting LED states
 - emergency fallback steps
-

15) Event Day Runbook (High-Level)

1. Place enclosure, plug in, flip master switch.
 2. Wait for master LED = Green.
 3. On admin device: open [event.local/admin](#).
 4. Create/select event; configure categories/weights.
 5. Start registration; print team codes.
 6. Start turn-in; scan submissions and teams.
 7. Lock turn-in; generate table assignment plan.
 8. Start judging:
 - table QR posters displayed
 - judges scan and enter seat
 9. Monitor completion dashboard; resolve stuck seats.
 10. Finalize scoring; publish results; generate reports.
 11. Export full event archive (CSV + PDFs) to external storage.
-